Part I: Pen and paper

1. **Compute the recall of a distance-weighted kNN with $k = 5$ and distance $d(x_1, x_2) =$ Hamming$(x_1, x_2) + \frac{1}{2}$ using leave-one-out evaluation schema (i.e., when classifying one observation, use all remaining ones).**

    For starters, it's worth noting that, in this context, the **Hamming distance** between two observations x_1 and x_2 is defined as the number of attributes that differ between them.

    Knowing this, we can now create an 8×8 matrix (as can be seen below), where each entry represents the Hamming distance $(+\frac{1}{2})$ between two observations. This matrix is symmetric, of course. Each column i, here, will have $8 - 1 = 7$ associated entries, each representing the distance d between the observation x_i and the remaining 7 observations: we will, then, pick the $k = 5$ nearest neighbors according to said distance, classifying x_i in a **distance-weighted** manner.

	x_1	x_2	x_3	x_4	x_5	x_6	x_7	x_8
x_1	×	5/2	3/2	1/2	3/2	3/2	3/2	5/2
x_2	5/2	×	3/2	5/2	3/2	3/2	3/2	1/2
x_3	3/2	3/2	×	3/2	5/2	5/2	1/2	3/2
x_4	1/2	5/2	3/2	×	3/2	3/2	3/2	5/2
x_5	3/2	3/2	5/2	3/2	×	1/2	5/2	3/2
x_6	3/2	3/2	5/2	3/2	1/2	×	5/2	3/2
x_7	3/2	3/2	1/2	3/2	5/2	5/2	×	3/2
x_8	5/2	1/2	3/2	5/2	3/2	3/2	3/2	×

Table 1: Distance d between observations - in teal, a given observation's k nearest neighbors

For each observation x_i, the k nearest neighbors are, therefore, the ones represented in teal in the table above. Instead of predicting a class given the neighbors' mode, we'll want to choose it in a distance-weighted manner here: this means that, for each observation x_i, we'll want to compute the **weighted majority vote** of its k nearest neighbors, where the weight of each neighbor is given by:

$$w_{i,j} = \frac{1}{d(x_i, x_j)}$$

where x_j is one of the k nearest neighbors of x_i. Considering the data gathered up until now, we'll have the following conclusions regarding the model's classification for each instance:

	Weighted distance to N	Weighted distance to P	Predicted Class
x_1	$3 * 1/(3/2) = 2$	$1/(3/2) + 1/(1/2) = 8/3$	P
x_2	$3 * 1/(3/2) + 1/(1/2) = 4$	$1/(3/2) = 2/3$	N
x_3	$1/(3/2) + 1/(1/2) = 8/3$	$3 * 1/(3/2) = 2$	N
x_4	$3 * 1/(3/2) = 2$	$1/(3/2) + 1/(1/2) = 8/3$	P
x_5	$1/(3/2) + 1/(1/2) = 8/3$	$3 * 1/(3/2) = 2$	N
x_6	$1/(3/2) + 1/(1/2) = 8/3$	$3 * 1/(3/2) = 2$	N
x_7	$2 * 1/(3/2) = 1/3$	$3 * 1/(3/2) + 1/(1/2) = 4$	P
x_8	$3 * 1/(3/2) = 2$	$1/(3/2) + 1/(1/2) = 8/3$	P

Table 2: Distance weighting for each observation

We'll have, given the data gathered above, the following confusion matrix:

Figure 1: Confusion Matrix

Moreover, the **recall** of a classifier is defined as the ratio between the number of true positives and the number of true positives plus the number of false negatives that the classifier makes. Looking at the confusion matrix above, we can assert that the associated recall will, therefore, be:

$$\frac{TP}{TP + FN} = \frac{2}{2 + 2} = \frac{2}{4} = 0.5$$

2. **Considering the nine training observations, learn a Bayesian classifier assuming: i) y_1 and y_2 are dependent, ii) $\{y_1, y_2\}$ and $\{y_3\}$ variable sets are independent and equally important, and iii) y_3 is normally distributed. Show all parameters.**

Considering both variable sets, $\{y_1, y_2\}$ and $\{y_3\}$, to be independent and equally important, it'll make sense to train a Naive Bayes classifier here, such that (and utilizing Bayes' theorem):

$$P(C =^N_P | y_1, y_2, y_3) = \frac{P(y_1, y_2, y_3 | C =^N_P)P(C =^N_P)}{P(y_1, y_2, y_3)}$$

More so, since $\{y_1, y_2\}$ and $\{y_3\}$ are independent (and y_1 and y_2 are dependent), we can rewrite the above as:

$$P(C =^N_P | y_1, y_2, y_3) = \frac{P(y_1, y_2 | C =^N_P)P(y_3 | C =^N_P)P(C =^N_P)}{P(y_1, y_2)P(y_3)}$$

According to the Naive Bayes' assumption, "the presence (or absence) of a particular feature in a class is unrelated to the presence (or absence) of any other feature", so we will, in fact, only need the above equation's numerator to be able to classify a new observation. The goal here is to find:

$$\text{argmax}_{c \in \{N, P\}} P(y_1, y_2 | C = c)P(y_3 | C = c)P(C = c)$$

For starters, we can note that, from the given training set:

$$P(C = P) = \frac{5}{9}, \quad P(C = N) = \frac{4}{9}$$

We also know that y_3 is normally distributed, meaning we'll have:

$$P(y_3 | C =^N_P) \sim \mathcal{N}(x | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

By also looking at the given training set (which includes the new, ninth sample), we'll be able to extrapolate the following probabilities:

$$C = N : P(C = N) = \tfrac{4}{9} \qquad\qquad C = P : P(C = P) = \tfrac{5}{9}$$

- $P(y_2 = A, y_2 = 0 | C = N) = 0$
- $P(y_2 = A, y_2 = 1 | C = N) = \tfrac{1}{4}$
- $P(y_1 = B, y_2 = 0 | C = N) = \tfrac{2}{4} = \tfrac{1}{2}$
- $P(y_1 = B, y_2 = 1 | C = N) = \tfrac{1}{4}$

- $P(y_1 = A, y_2 = 0 | C = P) = \tfrac{2}{5}$
- $P(y_2 = A, y_2 = 1 | C = P) = \tfrac{1}{5}$
- $P(y_1 = B, y_2 = 0 | C = P) = \tfrac{1}{5}$
- $P(y_1 = B, y_2 = 1 | C = P) = \tfrac{1}{5}$

Regarding y_3 and N labeled observations, we'll have the following parameters:

Regarding y_3 and P labeled observations, we'll have the following parameters:

$$\mu = \frac{1 + 0.9 + 1.2 + 0.8}{4} = 0.975$$

$$\mu = \frac{1.2 + 0.8 + 0.5 + 0.9 + 0.8}{5} = 0.84$$

$$\sigma^2 = \frac{1}{4 - 1} \sum_{i=1}^{4} (y_{3,i} - \mu)^2 = 0.029$$

$$\sigma^2 = \frac{1}{5 - 1} \sum_{i=1}^{5} (y_{3,i} - \mu)^2 = 0.063$$

The model is now ready to be used to classify new observations. Applying the model to the nine training observations, we'll have the following results:

	y_1	y_2	y_3	Class	$P(C = N)P(y_1, y_2, y_3 \mid C = N)$	$P(C = P)P(y_1, y_2, y_3 \mid C = P)$	Predicted Class	Veredict
x_1	A	0	1.2	P	0	$\tfrac{5}{9} \times \tfrac{2}{5} \times 0.92425$	P	TP
x_2	B	1	0.8	P	$\tfrac{4}{9} \times \tfrac{1}{4} \times 0.15206$	$\tfrac{5}{9} \times \tfrac{1}{5} \times 0.43669$	P	TP
x_3	A	1	0.5	P	$\tfrac{4}{9} \times \tfrac{1}{4} \times 0.00264$	$\tfrac{5}{9} \times \tfrac{1}{5} \times 0.08777$	N	FN
x_4	A	0	0.9	P	0	$\tfrac{5}{9} \times \tfrac{2}{5} \times 0.59446$	P	TP
x_5	B	0	1	N	$\tfrac{4}{9} \times \tfrac{1}{2} \times 0.55836$	$\tfrac{5}{9} \times \tfrac{1}{5} \times 0.73809$	N	TN
x_6	A	0	0.9	N	$\tfrac{4}{9} \times \tfrac{1}{2} \times 0.32981$	$\tfrac{5}{9} \times \tfrac{1}{5} \times 0.59446$	N	TN
x_7	B	1	1.2	N	$\tfrac{4}{9} \times \tfrac{1}{4} \times 0.9068$	$\tfrac{5}{9} \times \tfrac{1}{5} \times 0.92425$	P	FP
x_8	B	1	0.8	N	$\tfrac{4}{9} \times \tfrac{1}{4} \times 0.15206$	$\tfrac{5}{9} \times \tfrac{1}{5} \times 0.43669$	P	FP
x_9	B	0	0.8	P	$\tfrac{4}{9} \times \tfrac{1}{2} \times 0.15206$	$\tfrac{5}{9} \times \tfrac{1}{5} \times 0.43669$	P	FP

Table 3: Classification vs Real results

3. **Under a MAP assumption, compute** $P(Positive|x)$ **of each testing observation.**

Considering Bayes' theorem, we know $P(Positive|x)$ can also be written as follows:

$$P(Positive|x) = \frac{P(x|Positive)P(Positive)}{P(x)}$$

Since x has two independent variable sets, $\{y_1, y_2\}$ and $\{y_3\}$, we'll be able to write the above expression as follows (note that x's y_1 value will be written as y_1, and so on, for simplicity's sake):

$$P(Positive|x) = \frac{P(y_1, y_2|Positive)P(y_3|Positive)P(Positive)}{P(y_1, y_2)P(y_3)}$$

Let's consider the three samples given in the question's statement:

$$x_1' = \begin{pmatrix} A \\ 1 \\ 0.8 \end{pmatrix}, \quad x_2' = \begin{pmatrix} B \\ 1 \\ 1 \end{pmatrix}, \quad x_3' = \begin{pmatrix} B \\ 0 \\ 0.9 \end{pmatrix}$$

Considering the training observations, we can gather that there are:

x_1'	x_2'	x_3'
Among all (five) positive training samples, one with features $y_1 = A, y_2 = 1$, two with feature $y_3 = 0.8$. Among all (nine) training samples, two with features $y_1 = A, y_2 = 1$, three with feature $y_3 = 0.8$.	Among all (five) positive training samples, one with features $y_1 = B, y_2 = 1$, zero with feature $y_3 = 1$. Among all (nine) training samples, two with features $y_1 = B, y_2 = 1$, one with feature $y_3 = 1$.	Among all (five) positive training samples, one with features $y_1 = B, y_2 = 0$, one with feature $y_3 = 0.9$. Among all (nine) training samples, three with features $y_1 = B, y_2 = 0$, two with feature $y_3 = 0.9$.

Note, of course, that the probabilities regarding y_1 and y_2 features may be calculated in a discrete manner, directly utilizing the values gathered above. The same can't be said for y_3 with it being normally distributed: we'll need to calculate the parameters for it (note that Positive parameters were computed in the previous question), and only then will we be able to calculate the needed probabilities regarding these samples' y_3 values:

$$\mu_{y_3, Positive} = 0.84, \quad \mu_{y_3, All} = \frac{(1.2 + 0.8 + \cdots + 0.8)}{9} = 0.9 \qquad (1)$$

$$\sigma^2_{y_3, Positive} = 0.063, \quad \sigma^2_{y_3, All} = \frac{1}{9-1} \sum_{i=1}^{9} (y_{3,i} - \mu)^2 = 0.0475 \qquad (2)$$

x'_1	x'_2	x'_3
$P(y_3 = 0.8\|P) = 0.43669,$	$P(y_3 = 1\|P) = 0.73809,$	$P(y_3 = 0.9\|P) = 0.59446,$
$P(y_3 = 0.8) = 0.32314$	$P(y_3 = 1) = 0.67686$	$P(y_3 = 0.9) = 0.5$

We can, therefore, assert that:

$$P(Positive|x'_1) = \frac{1/5 \times 0.43669 \times 5/9}{2/9 \times 0.32314} = 0.6756978$$

$$P(Positive|x'_2) = \frac{1/5 \times 0.73809 \times 5/9}{2/9 \times 0.67686} = 0.5452309$$

$$P(Positive|x'_3) = \frac{1/5 \times 0.59446 \times 5/9}{3/9 \times 0.5} = 0.3963067$$

4. **Given a binary class variable, the default decision threshold of $\theta = 0.5$,**

$$f(x|\theta) = \begin{cases} Positive & \textbf{if } P(positive|x) > \theta \\ Negative & otherwise \end{cases}$$

can be adjusted. Which decision threshold – 0.3, 0.5 or 0.7 – optimizes testing accuracy?

As given by the question's statement, we know the actual class values of the testing observations. Moreover, we've calculated $P(Positive|x)$ for each of them in the previous question's answer, so we can easily calculate the accuracy of the classifier for each of the three decision thresholds (considering only these three testing samples):

	Class	$P(Positive|x)$	$\theta = 0.3$	$\theta = 0.5$	$\theta = 0.7$
x_1'	Positive	0.6756978	Positive	Positive	Negative
x_2'	Positive	0.5452309	Positive	Positve	Negative
x_3'	Negative	0.3963067	Positive	Negative	Negative

Table 4: $f(x|\theta)$ for varying thresholds, given three testing samples

Here, the classifier's accuracy, considering the table above, amounts to:

$$\text{Accuracy}(\theta = 0.3) = {}^{2}/_{3}$$
$$\text{Accuracy}(\theta = 0.5) = 1$$
$$\text{Accuracy}(\theta = 0.7) = {}^{1}/_{3}$$

As can be seen above, a threshold of $\theta = 0.5$ yields a higher accuracy than both 0.3 and 0.7, hence it should be chosen in order to optimize testing accuracy.

The code utilized in order to answer this section's first two questions answers can be found in this report's appendix. There's both a "general" code section, which can be found initially, and two functions, `first` and `second`, utilized to answer the respective questions - the first to plot the confusion matrices, while the second to test the given hypothesis.

5. **Using `sklearn`, considering a 10-fold stratified cross validation (`random=0`), plot the cumulative testing confusion matrices of kNN (uniform weights, $k = 5$, Euclidean distance) and Naïve Bayes (Gaussian assumption). Use all remaining classifier parameters as default.**

   We ended up utilizing `seaborn`'s `heatmap` function to plot the required confusion matrices side-by-side.


Figure 2: Cumulative testing confusion matrices of kNN and Naïve Bayes

6. **Using `scipy`, test the hypothesis "kNN is statistically superior to Naïve Bayes regarding accuracy", asserting whether is true.**

   `scipy.stats` provides the `ttest_rel` function, which can be utilized to test a hypothesis given two related samples - here, both kNN's and Naïve Bayes' scores are inherently related, since we're working with the same dataset on both classifiers, hence we can use this function to test the hypothesis.

   We opted to utilize it with a `alternative = 'greater'` parameter, since we're interested in testing whether kNN's accuracy is statistically superior to Naïve Bayes' accuracy, and not the other way around - as such, a "right-tailed" test is more appropriate. The considered hypotheses are, therefore:

   - H_0: kNN's accuracy is statistically equal to Naïve Bayes'
   - H_1: kNN's accuracy is statistically superior to Naïve Bayes'

   As a side-note, we've considered, in absence of a given confidence level in the question's statement, a confidence level of $1 - \alpha = 0.95$. After performing the test, we obtained a *p-value* of ≈ 0.9104 and a *t-statistic* of ≈ -1.457, which leads us to assert that, given $\alpha = 0.05$, we cannot reject the null hypothesis: we can't, therefore, assert whether kNN's accuracy is statistically superior to Naïve Bayes'.

7. **Enumerate three possible reasons that could underlie the observed differences in predictive accuracy between kNN and Naïve Bayes.**

   Note: the homework's FAQ states that the answer could state as low as two reasons, even though the original question's statement states three. We opted to answer the question as stated in the FAQ.

   TODO: add text

Appendix

```python
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import StratifiedKFold
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from scipy.io.arff import loadarff
from scipy.stats import ttest_rel

# Load data
data = loadarff('data/pd_speech.arff')
df = pd.DataFrame(data[0])
df['class'] = df['class'].str.decode('utf-8')

X = df.drop('class', axis=1)
y = df['class']

# Create 10-fold stratified cross validation
skf = StratifiedKFold(n_splits=10, shuffle=True, random_state=0)

# Create classifiers
knn = KNeighborsClassifier(n_neighbors=5, weights='uniform', metric='euclidean')
nb = GaussianNB()

knn_cm = np.zeros((2, 2))
nb_cm = np.zeros((2, 2))

knn_scores = np.array([])
nb_scores = np.array([])

for train_index, test_index in skf.split(X, y):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    knn.fit(X_train, y_train)
    nb.fit(X_train, y_train)

    knn_pred = knn.predict(X_test)
    nb_pred = nb.predict(X_test)

    # used in the first question
    knn_cm += confusion_matrix(y_test, knn_pred)
    nb_cm += confusion_matrix(y_test, nb_pred)

    # used in the second question
    knn_scores = np.append(knn_scores, accuracy_score(y_test, knn_pred))
    nb_scores = np.append(nb_scores, accuracy_score(y_test, nb_pred))

def first():
    """Used to answer the first question"""
    def create_plot(axis, labels, title, cm):
```

```python
55        plot = sns.heatmap(cm, annot=labels, cmap='Blues', ax=axis, fmt='s')
56        plot.set_title(title)
57        plot.set_ylabel('Predicted')
58        plot.set_xlabel('Real')
59        return plot
60
61    def label_heatmap(cm):
62        return np.asarray(
63            [f"{name}\n{count}" for name, count in zip(GROUP_NAMES, cm.flatten())]
64        ).reshape(2, 2)
65
66    GROUP_NAMES =['TN', 'FP', 'FN', 'TP']
67    KNN_LABELS = label_heatmap(knn_cm)
68    NB_LABELS = label_heatmap(nb_cm)
69    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(10, 5))
70    p1 = create_plot(ax1, KNN_LABELS, 'KNN Confusion Matrix', knn_cm)
71    p2 = create_plot(ax2, NB_LABELS, '(Gaussian) Naive Bayes Confusion Matrix', nb_cm
       )
72    plt.show()
73    # save the plot to a file
74    fig.savefig('assets/hw2-2.1.png')
75
76 def second():
77    """Used to answer the second question"""
78    return ttest_rel(knn_scores, nb_scores, alternative='greater')
79
80 # First programming question
81 first()
82
83 # Second programming question
84 results = second()
85 print("KNN better than NB?\np-value = {}\nt-statistic = {}".format(results.pvalue,
       results.statistic))
```